

Game theory

a course for the
MSc in ICT for Internet and multimedia

Leonardo Badia

leonardo.badia@gmail.com



Bayesian Games

From common knowledge to “just an idea”

Bayesian games

- Equilibrium concepts were developed earlier as modeling players that are forming both **consistent** and **correct** beliefs
 - ▣ e.g., the payoff of other players must be known
- Harsanyi (1960) proposed that this can be incorporated into the concept of beliefs
 - ▣ We had beliefs about other players' moves (with certainty over their preferences and costs)
 - ▣ Now, we include beliefs about these characteristics as well!

Bayesian games

- Game of **incomplete information**: beliefs over the characteristics of other players are captured by their **types**
- Players can be of different types, which implies them to behave differently and also the other players to have beliefs about it
- We will develop an equilibrium concept that still requires these beliefs to be consistent and correct (so we use the same procedures)

Preliminary Nature's move


- Players may have different payoffs (each represented by a type)
- The timing of the game is as follows
 - ▣ Nature draws a type vector (t_1, \dots, t_n) among all possibilities for all the players
 - ▣ Nature reveals type t_i to player i only
 - ▣ Players choose their actions
 - ▣ Payoffs are computed
- This is a dynamic game where the players do not know Nature's move in its entirety

Remark on types

- Actually, types can represent more than just different preferences
 - ▣ Types may not differ in the player's preferences but in the **knowledge** that a player has about the types of the other players or some other characteristics of the game
 - ▣ More subtle variant: we keep it for a later stage

Example: Entry Game

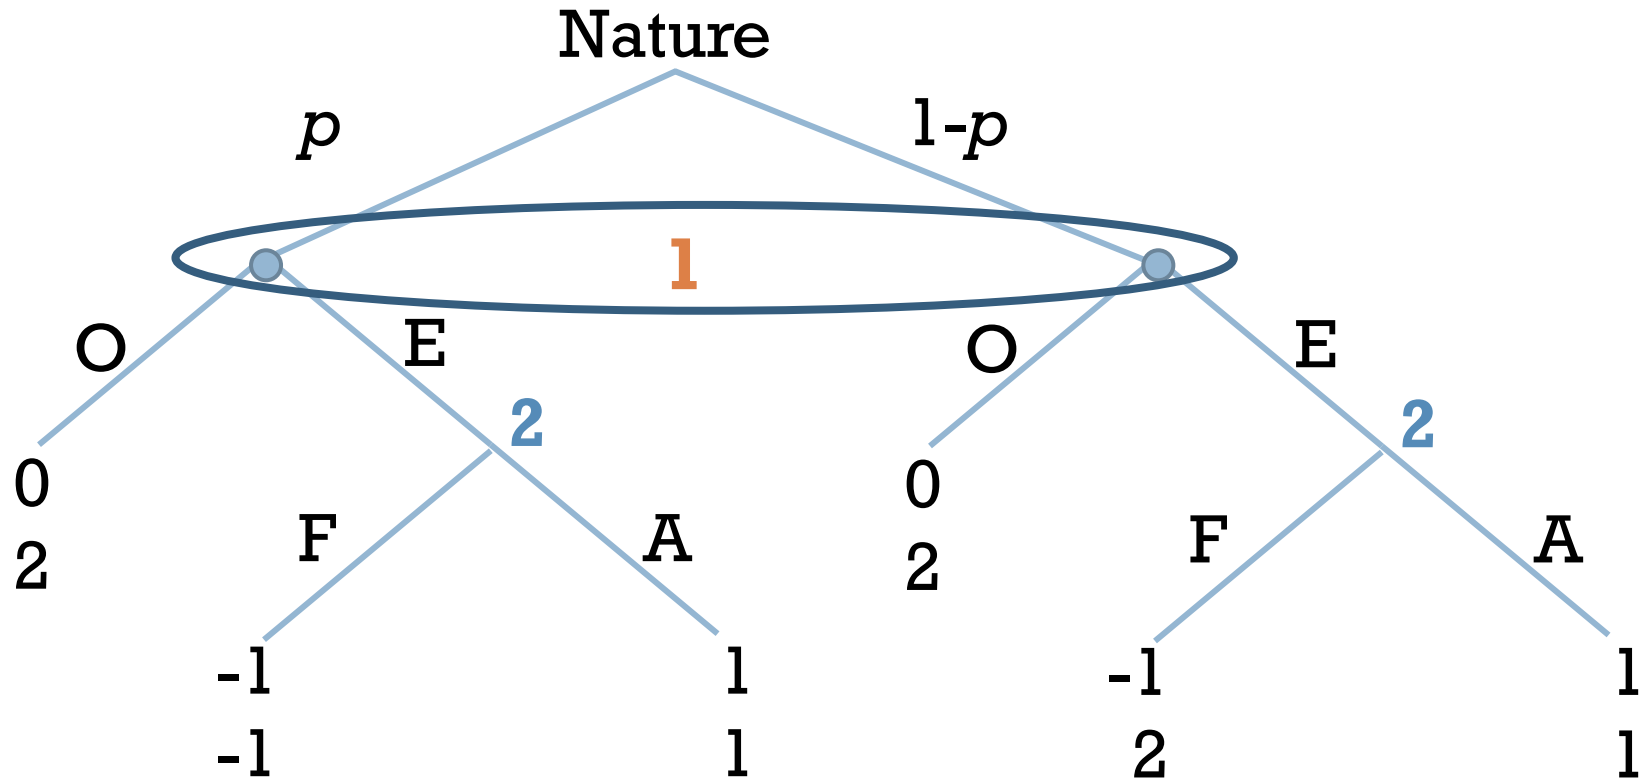
- Player 1 is a newcomer (e.g., in a market or network); he may (E)nter or stay (O)ut
- Player 2 is incumbent, if 1 enters, 2 may (A)ccept or (F)ight
- SPE outcome is (E,A)
 - ▣ there is another NE outcome that is not SPE and it is (O,F); check it!



Example: Entry Game

- Now, assume player 2 can be of two types
 - ▣ Rational: behaves as already discussed
 - ▣ Crazy: enjoys fighting and his/her payoff for (Enter, Fight) is actually 2 instead of -1
- Nature sets the type of 2: with probability p to be “Rational” (with $1-p$ he/she is “Crazy”)
- Player 1 does not know the type of 2, and has non-singleton information set about it
 - ▣ Player 2 instead knows his/her own type

Entry game, with types



Entry game, beliefs

- Players know their type, not the opponents'
- How can they well form a best response?
 - ▣ They need to create beliefs about these types
 - ▣ For this to be possible, we assume that they do not precisely know the types of their opponents, but know how they are set
 - ▣ Thus, they know the **probability distribution** of the opponents' types: it is common knowledge!
 - ▣ This is called **common prior assumption**

Entry game, strategies

- Also strategies need to be expanded
 - ▣ Apparently, player **2** does not have much choice even though he/she moves only if **1** chooses E
- In normal form, **2** has **four** pure strategies:
two actions \times two information sets
 - ▣ Because of Nature's move, **2** can have two types (**1**'s move is irrelevant instead: check yourself)
- In other words, **2**'s strategy is (xy) where:
 - ▣ x describes what a rational player **2** does
 - ▣ y describes what a crazy player **2** does

Bayesian strategies

- In incomplete information games, strategies are further expanded by stating what each type of a player does
 - ▣ The number of combinations explodes soon!
- In our example of Entry game:
 - ▣ player 1 has two strategies (O,E)
 - ▣ player 2 has four strategies (AA, AF, FA, FF)
- Each pair of pure strategies sets the path of play, which also depends on Nature's choice

Bayesian game, normal form

- For example, if the game is played as (E, AF)
 - ▣ player 1 gets: $p \times 1 + (1-p) \times (-1) = 2p-1$
 - ▣ player 2 gets: $p \times 1 + (1-p) \times 2 = 2-p$

		player 2			
		AA	AF	FA	FF
player 1	O	0, 2	0, 2	0, 2	0, 2
	E	1, 1	$2p-1, 2-p$	$1-2p, 1-2p$	$-1, 2-3p$

Bayesian game, normal form

- The value of p matters to determine the NEs
- For example, set $p = \frac{2}{3}$

		player 2			
		AA	AF	FA	FF
player 1	O	0, 2	0, 2	0, 2	0, 2
	E	1, 1	$\frac{1}{3}, \frac{4}{3}$	$-\frac{1}{3}, -\frac{1}{3}$	-1, 0

- If we are agnostic on the matrix and just seek NEs, we found three: (O,FA), (O,FF), and (E,AF)

Bayesian games, discussion

- Actually the reasoning can be extended to mixed strategies in a straightforward way
- The reasoning still applies expected payoffs
 - ▣ but now derived from the probability of **types**!
- We change incomplete information into imperfect information
 - ▣ Specifically, uncertainty is now on Nature's choice of players' types at the beginning

Bayesian games, discussion

- In this example, player 1 just has one type
 - ▣ Clearly he wants to average on all possibilities of his opponent's (player 2) types
- For player 2, things are a bit more bizarre
 - ▣ He needs to take expectation on his own types!
 - ▣ Real players (Rational 2 or Crazy 2) are replaced by a meta-player MP2 playing for both sometimes called “type agent” representation
- Do the best responses of MP2 = reality?
 - ▣ Yes, due to equivalence normal → extensive form

Bayesian games, discussion

- Evaluating expectations means that the probability distribution of types (not the types themselves) is common knowledge
- This is a strong assumption that is required to compute equilibria, because players must be able to conjecture on the game outcome
- The realism of this assumption may be critical and needs to be checked every time

Bayesian game: definition

Formal definition of the terms

Representing Bayesian games

- Our formalization of games used up to now a normal-form representation including:
 - ▣ set of players $\mathcal{N} = 1, \dots, n$
 - ▣ strategy space of each player S_i (for $i = 1..n$)
 - ▣ utilities of players $u_i : (S_1, S_2, \dots, S_n) \rightarrow \mathbb{R}$ (for $i = 1..n$)
- Bayesian games add three more ingredients:
 - ① **type & type space** of each player T_i (for $i = 1..n$)
 - ② also, utilities are now **type-dependent**
 - ③ finally, we need **beliefs** about other players types

Static Bayesian game

- We consider a static game: n players, each player's strategy is just an action a_i in set A_i
- ① Player i 's type is $t_i \in T_i$, chosen by Nature for each player from 1 to n through the **prior** probability distribution $\phi(t_1, \dots, t_n)$
 - the prior is **common knowledge** among players
- ② About the influence on types over payoffs, we make the assumption of **private values**
 - This means that $u_i = u_i(a_1, a_2, \dots, a_n, t_i)$
 - If $u_i = u_i(a_1, a_2, \dots, a_n, t_1, t_2, \dots, t_n)$: **common values**

Type of a player

- For example, player i can have two different payoff functions $u_{i,a}(a_i, a_{-i})$, $u_{i,b}(a_i, a_{-i})$
 - we represent this by setting type space $T_i = \{t_a, t_b\}$ and imposing $u_i(a_i, a_{-i}; t_j) = u_{i,j}(a_i, a_{-i})$
- Types can be used to limit available actions
 - If a player has actions $\{F, G, H\}$, but H is permitted only with probability q , we define types t_a and t_b
 - t_a and t_b have respective probabilities q and $1 - q$
- In both cases $\{F, G, H\}$ are feasible actions, but all payoffs of move H under type t_b are $-\infty$

Beliefs on types

- ③ Because players know their own types, they can form **beliefs** on the other opponent by simply applying conditional probability:

$$\phi(t_{-i} | t_i) = \phi(t_1, \dots, t_n) / \phi(t_i)$$

because players know their own types!

- **Comment.** Actually, we should include the entire hierarchy of “beliefs about other players’ beliefs” that soon gets complicated
 - It can be shown that it is equivalent to “merge” down the hierarchy just to 1st-order beliefs

Beliefs on types (comments)

- Types are correlated; they are independent if
$$\phi(t_1, \dots, t_n) = \phi(t_1) \cdot \dots \cdot \phi(t_n)$$
- *i* knows its own type, but not others' (t_{-i});
he estimates them via **belief** $\phi_i(t_{-i} | t_i)$
 - ▣ Prior versus posterior probabilities
- Our assumption of the prior being common knowledge equals to **perfect information**
 - ▣ In the case of incomplete and imperfect info.,
belief $\phi_i(t_{-i} | t_i)$ may even be wrong and have
nothing to do with the true prior

Static Bayesian game

- Static Bayesian needs
 - players ^①
 - action spaces ^②
 - beliefs ^④
 - type spaces ^③
 - (type-dep.) payoffs ^⑤
- $G = \{\mathcal{N}^{①}; A_1, \dots, A_n^{②}; T_1, \dots, T_n^{③}; \phi_1, \dots, \phi_n^{④}; u_1, \dots, u_n^{⑤}\}$
 - where $u_i = u_i(a_1, \dots, a_n; t_i)$.
- A pure strategy for i is a map $s_i: T_i \rightarrow A_i$, i.e., it tells what i plays as his/her type is known
- A mixed strategy for i is a probability distribution over pure strategies

Strategies of Bayesian games

- Type-contingent definition of pure/mixed strategies → reminiscent of dynamic games
 - ▣ We can think of a general strategy as being defined before the type of i is even set!
 - ▣ Player i decides a strategy $s_i : T_i \rightarrow A_i$:
then, if her type is $t_i \in T_i$, she will play $s_i(t_i)$
 - ▣ This is useful because it allows **other players** to create beliefs over the strategy of a player i who can be of different types

Strategies of Bayesian games

take $p = 2/3$

		player 2			
		AA	AF	FA	FF
player 1	O	0, 2	0, 2	0, 2	0, 2
	E	1, 1	$1/3, 4/3$	$-1/3, -1/3$	-1, 0

- 2's strategies depend on Nature setting $t_2 \in \{r, c\}$
- If 1 believes 2 uses pure strategy $s_2 = AF$: $\begin{cases} s_2(r) = A \\ s_2(c) = F \end{cases}$
- 1's expected payoff when playing E is

$$\mathbb{E}[u_1(E, s_2)] = p u_1(E, s_2(r)) + (1-p) u_1(E, s_2(c))$$
 that for $p = 2/3$ is exactly the matrix entry for (E, AF)

Strategies of Bayesian games

- **Comment 1.** Assume a typed player i is using a type-dependent pure strategy, and Nature randomly chooses his/her type t_i
 - ▣ from the point of view of the opponents $-i$, they are facing a player using mixed strategies
 - ▣ in both cases an expectation is taken: for the opponents $-i$ it does not matter if it is i 's own decision to randomize the action or just a consequence of its (Nature-chosen) type

Strategies of Bayesian games

- **Comment 2.** When specifying a strategy of a typed player, must we be type-dependent?
 - ▣ Why specify what i will do if Nature chooses any of i 's types, if in the end only one is real?
- No need for i itself.. but for the opponents $-i$!
 - ▣ This way, they can form beliefs over i 's behavior
 - ▣ And, combined with the posterior (giving the condition on i 's type), compute expected payoffs
- We did a similar thing for dynamic games!

Bayesian Nash equilibrium

- It is a Nash equilibrium in Bayesian games
- In $G = \{\mathcal{N}; A_1, \dots, A_n; T_1, \dots, T_n; \phi_1, \dots, \phi_n; u_1, \dots, u_n\}$, joint strategy $s^* = (s_1^*, \dots, s_n^*)$ is said to be a **Bayesian Nash Equilibrium** if for each player i and each type $t_i \in T_i$, $s_i(t_i)$ maximizes the expected payoff.

That is

$$\max_{s_i \in S_i} \sum_{t_{-i}} u_i(s_1^*(t_1), \dots, s_{i-1}^*(t_{i-1}), s_i, s_{i+1}^*(t_{i+1}), \dots, s_n^*(t_n), t_i) \phi_i(t_{-i} | t_i)$$

Bayesian Nash equilibrium

- This can be rewritten as:

$$\mathbb{E}[u_i(s_i^*(t_i), s_{-i}^*(t_{-i}), t_i) | t_i] \geq \mathbb{E}[u_i(s_i, s_{-i}^*(t_{-i}), t_i) | t_i]$$

for every $s_i \in S_i$

- In other words, i does not want to change strategy given the information he/she knows
 - ▣ strategy of i = a choice of action for each type
 - ▣ what i does not know, he just estimates!
- This definition can be generalized according to the type space (if continuous \rightarrow integrals)

mixed strategies / types

Harsanyi's interpretation of beliefs

back to Battle of the Sexes

		Brian	
		R	S
Ann	R	2, 1	0, 0
	S	0, 0	1, 2

- Mixed strategy = probabilities to play R
- NEs of the game: (0,0), (1,1) and $(\frac{2}{3}, \frac{1}{3})$
- The mixed NE can be seen as **pure** BNE of a related game with a bit of incomplete info

back to Battle of the Sexes

		Brian	
		R	S
Ann	R	$2 + \mathbf{c}, 1$	$0, 0$
	S	$0, 0$	$1, 2 + \mathbf{d}$

- Ann and Brian do not know each other well
- Increase Ann's payoffs at (R,R) and Brian's at (S,S) by \mathbf{c} and \mathbf{d} , both (\mathbf{c} and \mathbf{d}) falling in $[0, x]$
 - ▣ Think of x as a “perturbation”
 - ▣ The exact amount of \mathbf{c} (or \mathbf{d}) is privately known by Ann (or Brian) only: type of the player

back to Battle of the Sexes

- Ann's strategy: play R if $c > C$, otherwise play S
 - Same for Brian, choose S if $d > D$, else play R
- This strategy is in fact a Bayesian NE
- Ann does not know d ! Her expected payoff is
$$D/x (2+c) + (1 - D/x) \cdot 0 = (2+c) D/x$$
if she plays R
$$D/x \cdot 0 + (1 - D/x) \cdot 1 = 1 - D/x$$
if she plays S
- Thus, Ann plays R if $c \geq \boxed{x/D - 3} = C$

back to Battle of the Sexes

- We have $x / \mathbf{D} - 3 = \mathbf{C}$
- Similarly, Brian's expected payoff is
 $(1 - \mathbf{C}/x) \cdot 0 + \mathbf{C}/x (2 + \mathbf{d}) = (2 + \mathbf{d}) \mathbf{C}/x$
if he plays S
 $(1 - \mathbf{C}/x) \cdot 1 + \mathbf{C}/x (2 + \mathbf{d}) \cdot 0 = 1 - \mathbf{C}/x$
if he plays R
- Brian plays S if $\mathbf{d} \geq x / \mathbf{C} - 3$. Thus, $x / \mathbf{C} - 3 = \mathbf{D}$.
- Combining these two conditions we have
 $x / \mathbf{D} - 3 = \mathbf{C}$, $x / \mathbf{C} - 3 = \mathbf{D}$
Thus, $\mathbf{C} = \mathbf{D}$ and $\mathbf{C}^2 + 3\mathbf{C} - x = 0$

back to Battle of the Sexes

□ Solving $C^2 + 3C - x = 0$

□ $C = \frac{-3 + \sqrt{9+4x}}{2}$

□ The probability of playing R for Ann is thus $1-C/x$, i.e.,

$$\frac{2x + 3 - \sqrt{9+4x}}{2x} \xrightarrow{x \rightarrow 0} \frac{2}{3}$$